

```
In [21]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

In [3]: `pip install Seaborn`

## Collecting Seaborn

```

  Downloading seaborn-0.12.2-py3-none-any.whl (293 kB)
      0.0/293.3 kB ? eta -:--:--
      0.0/293.3 kB ? eta -:--:--
      -      10.2/293.3 kB ? eta -:--:--
      -----      41.0/293.3 kB 667.8 kB/s eta 0:0
0:01
      -----      51.2/293.3 kB 440.4 kB/s eta 0:0
0:01
      -----      71.7/293.3 kB 393.8 kB/s eta 0:0
0:01
      -----      71.7/293.3 kB 393.8 kB/s eta 0:0
0:01
      -----      102.4/293.3 kB 393.8 kB/s eta 0:0
0:01
      -----      122.9/293.3 kB 380.5 kB/s eta 0:0
0:01
      -----      122.9/293.3 kB 380.5 kB/s eta 0:0
0:01
      -----      174.1/293.3 kB 420.1 kB/s eta 0:0
0:01
      -----      184.3/293.3 kB 413.5 kB/s eta 0:0
0:01
      -----      204.8/293.3 kB 429.5 kB/s eta 0:0
0:01
      -----      225.3/293.3 kB 430.1 kB/s eta 0:0
0:01
      -----      256.0/293.3 kB 450.1 kB/s eta 0:0
0:01
      -----      293.3/293.3 kB 452.9 kB/s eta 0:0
0:00
Requirement already satisfied: numpy!=1.24.0,>=1.17 in c:\users\dell\appdata\local\programs\python\python311\lib\site-packages (from Seaborn) (1.24.3)
Requirement already satisfied: pandas>=0.25 in c:\users\dell\appdata\local\programs\python\python311\lib\site-packages (from Seaborn) (2.0.1)
Requirement already satisfied: matplotlib!=3.6.1,>=3.1 in c:\users\dell\appdata\local\programs\python\python311\lib\site-packages (from Seaborn) (3.7.1)
Requirement already satisfied: contourpy>=1.0.1 in c:\users\dell\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.1->Seaborn) (1.0.7)
Requirement already satisfied: cycler>=0.10 in c:\users\dell\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.1->Seaborn) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\dell\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.1->Seaborn) (4.39.4)
Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\dell\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.1->Seaborn) (1.4.4)
Requirement already satisfied: packaging>=20.0 in c:\users\dell\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.1->Seaborn) (23.1)
Requirement already satisfied: pillow>=6.2.0 in c:\users\dell\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.1->Seaborn) (9.5.0)
Requirement already satisfied: pyparsing>=2.3.1 in c:\users\dell\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.1

```

->Seaborn) (3.0.9)

Requirement already satisfied: python-dateutil>=2.7 in c:\users\dell\appdata\local\programs\python\python311\lib\site-packages (from matplotlib!=3.6.1,>=3.1->Seaborn) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in c:\users\dell\appdata\local\programs\python\python311\lib\site-packages (from pandas>=0.25->Seaborn) (2023.3)

Requirement already satisfied: tzdata>=2022.1 in c:\users\dell\appdata\local\programs\python\python311\lib\site-packages (from pandas>=0.25->Seaborn) (2023.3)

Requirement already satisfied: six>=1.5 in c:\users\dell\appdata\local\programs\python\python311\lib\site-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.1->Seaborn) (1.16.0)

Installing collected packages: Seaborn

Successfully installed Seaborn-0.12.2

Note: you may need to restart the kernel to use updated packages.

```
In [22]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [30]: df=pd.read_csv(r"D:\Users\DELL\Desktop\dodo\bottle.csv")
df
```

C:\Users\DELL\AppData\Local\Temp\ipykernel\_14204\2765619747.py:1: DtypeWarning: Columns (47,73) have mixed types. Specify dtype option on import or set low\_memory=False.

```
df=pd.read_csv(r"D:\Users\DELL\Desktop\dodo\bottle.csv")
```

Out[30]:

	Cst_Cnt	Btl_Cnt	Sta_ID	Depth_ID	Depthm	T_degC	Salnty	O2ml_L	STheta	O2%
<b>0</b>	1	1	054.0 056.0	19- 4903CR- HY-060- 0930- 05400560- 0000A-3	0	10.500	33.4400	NaN	25.64900	N
<b>1</b>	1	2	054.0 056.0	19- 4903CR- HY-060- 0930- 05400560- 0008A-3	8	10.460	33.4400	NaN	25.65600	N
<b>2</b>	1	3	054.0 056.0	19- 4903CR- HY-060- 0930- 05400560- 0010A-7	10	10.460	33.4370	NaN	25.65400	N
<b>3</b>	1	4	054.0 056.0	19- 4903CR- HY-060- 0930- 05400560- 0019A-3	19	10.450	33.4200	NaN	25.64300	N
<b>4</b>	1	5	054.0 056.0	19- 4903CR- HY-060- 0930- 05400560- 0020A-7	20	10.450	33.4210	NaN	25.64300	N
...	...	...	...	...	...	...	...	...	...	...
<b>864858</b>	34404	864859	093.4 026.4	20- 1611SR- MX-310- 2239- 09340264- 0000A-7	0	18.744	33.4083	5.805	23.87055	108
<b>864859</b>	34404	864860	093.4 026.4	20- 1611SR- MX-310- 2239- 09340264- 0002A-3	2	18.744	33.4083	5.805	23.87072	108
<b>864860</b>	34404	864861	093.4 026.4	20- 1611SR- MX-310- 2239- 09340264- 0005A-3	5	18.692	33.4150	5.796	23.88911	108
<b>864861</b>	34404	864862	093.4 026.4	20- 1611SR- MX-310- 2239- 09340264- 0010A-3	10	18.161	33.4062	5.816	24.01426	107

	Cst_Cnt	Btl_Cnt	Sta_ID	Depth_ID	Depthm	T_degC	Salnty	O2ml_L	STheta	O2%
864862	34404	864863	093.4 026.4	20- 1611SR- MX-310- 2239- 09340264- 0015A-3	15	17.533	33.3880	5.774	24.15297	105

864863 rows × 74 columns


```
In [31]: df=df[['Salnty','T_degC']]
df.columns=['Sal','Temp']
df.head(10)
```

Out[31]:

	Sal	Temp
0	33.440	10.50
1	33.440	10.46
2	33.437	10.46
3	33.420	10.45
4	33.421	10.45
5	33.431	10.45
6	33.440	10.45
7	33.424	10.24
8	33.420	10.06
9	33.494	9.86

```
In [32]: sns.lmplot(x="Sal",y="Temp",data=df,order=2,ci=None)
```

```
Out[32]: <seaborn.axisgrid.FacetGrid at 0x21303554310>
```



```
In [33]: df.describe()
```

```
Out[33]:
```

	Sal	Temp
<b>count</b>	817509.000000	853900.000000
<b>mean</b>	33.840350	10.799677
<b>std</b>	0.461843	4.243825
<b>min</b>	28.431000	1.440000
<b>25%</b>	33.488000	7.680000
<b>50%</b>	33.863000	10.060000
<b>75%</b>	34.196900	13.880000
<b>max</b>	37.034000	31.140000



In [34]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 864863 entries, 0 to 864862
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0    Sal      817509 non-null  float64
 1   Temp      853900 non-null  float64
dtypes: float64(2)
memory usage: 13.2 MB
```

In [66]: `df.fillna(method='ffill',inplace=True)`

C:\Users\DELL\AppData\Local\Temp\ipykernel\_14204\4116506308.py:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))  
`df.fillna(method='ffill',inplace=True)`

In [45]: `x=np.array(df['Sal']).reshape(-1,1)`  
`y=np.array(df['Temp']).reshape(-1,1)`  
`df.dropna(inplace=True)`


C:\Users\DELL\AppData\Local\Temp\ipykernel\_14204\1516682253.py:3: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))  
`df.dropna(inplace=True)`

In [50]: `x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.5)`  
`regr=LinearRegression()`  
`regr.fit(x_train,y_train)`  
`print(regr.score(x_test,y_test))`

0.25406974224573897


```
In [52]: y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```



```
In [65]: df500=df[:][:500]
```

```
In [55]: sns.lmplot(x="Sal",y="Temp",data=df500,order=1,ci=None)
```

```
Out[55]: <seaborn.axisgrid.FacetGrid at 0x213540b8690>
```



```
In [68]: df500.fillna(method='ffill',inplace=True)
x=np.array(df['Sal']).reshape(-1,1)
y=np.array(df['Temp']).reshape(-1,1)
df.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(x_train,y_train)
print("Regression:",regr.score(x_test,y_test))
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```


C:\Users\DELL\AppData\Local\Temp\ipykernel\_14204\574439154.py:4: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy) ([https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy))

```
df.dropna(inplace=True)
```

Regression: 0.2513136615661431



```
In [69]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2_score:",r2)
```

R2\_score: 0.2513136615661431

In [ ]: